
When Does a Math Model Know the Reasoning Went Wrong?

Tracing the First Error in Hidden States

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Language models can produce fluent mathematical solutions even after the rea-
2 soning has failed. We ask whether a math-specialized model’s hidden states reg-
3 ister the first invalid step. At every retained step boundary in the 3,400-trace Pro-
4 cessBench dataset, we record the residual-stream state of Qwen2.5-Math-1.5B-
5 Instruct and train layer-wise linear probes to predict whether the solution is in-
6 valid so far. With problem-grouped splits and validation-only model selection,
7 the selected probe obtains held-out AUROC 0.866 (95% whole-trace bootstrap CI
8 [0.849, 0.884]), compared with 0.733 for current-step TF-IDF and 0.730 for posi-
9 tion. The same score reaches AUROC 0.873 when restricted to erroneous traces,
10 and source-specific directions transfer across every pair of four mathematical do-
11 mains (mean off-diagonal AUROC 0.805). Temporal localization is harder: the
12 probe identifies the exact first error on 29.0% of erroneous traces, lands within one
13 step on 59.3%, and is systematically late by 0.59 steps among detections. Finally,
14 adding or subtracting the probe direction at the boundary does not change the
15 model’s explicit correctness verdict. Because that native verdict readout is itself
16 degenerate (AUROC 0.283), the intervention is inconclusive rather than evidence
17 of causal non-use. The hidden states contain a transferable invalidity signal, but
18 the experiment does not show precise timing or behavioral use.

19 1 Introduction

20 Outcome accuracy hides an important failure mode of mathematical reasoning: a solution can reach
21 a wrong answer through one early mistake, or reach the right answer through an invalid derivation.
22 Verifiers therefore increasingly score intermediate steps rather than only final answers (Cobbe et al.,
23 2021; Lightman et al., 2023). Yet an external verifier leaves open a mechanistic question: when a
24 solution first becomes invalid, does the solving model’s own representation change in a detectable
25 way?

26 An internal signal near the first invalid step would provide evidence that the model tracks errors
27 during an extended computation. It might also support lightweight detection. But three claims
28 must be separated. A state may contain decodable information without localizing the transition;
29 a direction may rank examples across domains without transferring its threshold; and the model’s
30 output computation need not use a decodable direction (Hewitt and Liang, 2019; Elazar et al., 2021).

31 We study these distinctions in Qwen2.5-Math-1.5B-Instruct (Yang et al., 2024) on ProcessBench,
32 whose human annotations identify the first erroneous step in solutions drawn from GSM8K, MATH,
33 OlympiadBench, and Omni-MATH (Zheng et al., 2025). Our experiment was specified before test
34 results were inspected. We extract one residual-stream state per layer and written step, fit regularized

linear probes on problem-grouped splits, evaluate first-threshold-crossing localization, test source-to-source transfer, and intervene along the learned direction.

Our main findings are:

- Invalid-so-far is linearly decodable at the selected hidden-state index with AUROC 0.866, exceeding position, current-step lexical, and shuffled-label controls. Discrimination remains high within erroneous traces (AUROC 0.873).
- The ranking signal transfers across all 12 off-diagonal source pairs, but exact localization is weaker and less calibrated: only 29.0% of erroneous traces are localized exactly.
- The preregistered additive intervention yields no signed dose response. However, the underlying explicit verdict score fails to distinguish valid from invalid boundaries, so this result cannot adjudicate whether the decoded signal is causally used.

2 Related work

Process-level mathematical verification. Sampling solutions and ranking them with a verifier improved GSM8K performance before modern process supervision (Cobbe et al., 2021). PRM800K then enabled direct comparison of outcome and human step-level supervision; process supervision performed better on a representative MATH subset (Lightman et al., 2023). Later step-aware verifiers and process reward models used intermediate correctness for search and training (Li et al., 2023b). These reward models can generalize poorly from GSM8K and MATH to harder competition problems. ProcessBench addresses this problem by annotating the earliest error in 3,400 traces spanning four sources and evaluating both reward and prompted critic models (Zheng et al., 2025). Prompted LLM studies likewise find that identifying and correcting mathematical errors is a separate capability from solving, and that exact first-error localization remains difficult even when a reference solution is provided (Li et al., 2024; Srivatsa et al., 2025).

We use ProcessBench differently from its standard benchmark setting. We repartition it to analyze one model’s representations, so our metrics are not leaderboard scores. Rather than score generic step quality, we ask whether a prefix has crossed the human-annotated first-error boundary and when the probe first detects that change.

Internal correctness and reasoning signals. Models can sometimes estimate whether their own answers are correct (Kadavath et al., 2022), and truth-related information can be recovered from activations even when outputs are unreliable (Burns et al., 2023). Work closest to ours shows that lightweight probes detect arithmetic errors and can guide selective re-prompting (Sun et al., 2025); ReProbe uses learned hidden-state scorers for multi-domain step verification (Ni et al., 2026); and STEP uses hidden states to prune low-quality reasoning traces during test-time scaling (Liang et al., 2026). Circuit analysis of simple arithmetic further finds a dissociation between computation and validation, with smaller models relying heavily on surface-level numerical consistency heads (Bertolazzi et al., 2025).

These studies show that internal states can predict reasoning quality. We use human first-error labels for open-ended model-generated solutions from grade-school through Olympiad mathematics. We measure the first score crossing rather than only step classification, test every source-to-source transfer, and evaluate the probe direction with a held-out causal intervention.

From probing to causal evidence. Probe accuracy alone is ambiguous because a classifier may exploit accessible but behaviorally unused information; controls and constrained probe families help but do not establish use (Hewitt and Liang, 2019; Belinkov, 2022). Causal work instead modifies representations through amnesic removal (Elazar et al., 2021), causal tracing (Meng et al., 2022), activation addition, or inference-time steering (Turner et al., 2023; Li et al., 2023a). We therefore treat intervention as a stronger, separate test. A null intervention is interpretable only when its output assay is sensitive and meaningful. Our original verdict prompt fails that condition.




Figure 1: Method and one held-out score trajectory. We probe boundary residual states at every layer, select all hyperparameters on validation data, and define the predicted first error as the first threshold crossing. In the example, the human error begins at step 1 and the score crosses at step 2. High step AUROC therefore need not imply exact localization.

3 Method

3.1 Data, model, and leakage-resistant splits

ProcessBench contains 3,400 model-generated solutions with a human label e_i : the zero-indexed first erroneous step, or -1 if every step is correct (Zheng et al., 2025). Its four sources span grade-school word problems, competition mathematics, and Olympiad problems (Cobbe et al., 2021; Hendrycks et al., 2021; He et al., 2024; Gao et al., 2024). We evaluate the 1.5B-parameter instruction-tuned Qwen2.5-Math model in bfloat16 with a 2,048-token context limit; over-length traces are excluded rather than truncated.

All traces sharing whitespace-normalized, case-folded problem text form one group. Within source $\times$ error-status strata, a deterministic seed-42 allocation places approximately 60%/20%/20% of groups in train/validation/test. All steps from a problem therefore remain in one partition. The published test artifact contains 669 traces (432 erroneous and 237 fully correct) and 4,985 step boundaries.

3.2 Boundary representations and targets

The prompt renders each step k followed by a unique END_STEP marker. One causal forward pass records the residual-stream state $h_{ik}^{(\ell)} \in \mathbb{R}^{1536}$ at each marker across 29 hidden-state indices. Although later steps occur in the same prompt, causal attention prevents them from influencing an earlier boundary state. For trace i and step k , the primary target is

$$y_{ik} = \mathbb{I}[e_i \geq 0 \wedge k \geq e_i], \quad (1)$$

which denotes that the displayed reasoning is *invalid so far*. It is intentionally persistent after onset and should not be read as a label of the current step in isolation.

3.3 Predictive, localization, and transfer tests

At each index ℓ , we standardize dimensions using training data and fit a class-balanced, L2-regularized logistic probe. Validation AUROC selects $C \in \{0.01, 0.1, 1, 10\}$; validation Process F1 selects a threshold from 0.05 to 0.95; and validation Process F1 (then AUROC) selects the layer. The chosen probe is refit on train+validation and evaluated once on test. The selected setting is index 23, $C = 0.01$, threshold 0.645.

For a trace, the predicted first error is the first score at or above this threshold, or -1 if none crosses. We report error-trace exact accuracy, correct-trace rejection accuracy, their harmonic mean (Process F1), complete-trace exact outcome, and tolerance-based localization. We also report step AUROC within erroneous traces, preventing success from a constant distinction between wholly correct and




Figure 2: Held-out decoding, localization, and controls. Error bars at the selected index are 95% whole-trace bootstrap intervals. The hidden-state probe ranks invalid prefixes well, but its first-error Process F1 is much lower than its step AUROC. “Within-error traces” has no meaningful correct-trace rejection term, so its Process F1 is shown as zero.

erroneous solutions. Controls use absolute/fractional position, current-step TF-IDF unigrams and bigrams, and shuffled training labels. Uncertainty uses 1,000 percentile bootstrap resamples of whole traces.

For transfer, a separate index-23 probe is trained on each source and evaluated on every test source. Its threshold is fixed using only the training source’s validation set. AUROC therefore tests ranking transfer, whereas Process F1 additionally tests probability and threshold transfer.

3.4 Probe-direction intervention

The standardized probe weight is transformed back to raw activation coordinates and normalized to v . At the input to decoder layer 23, we replace one held-out boundary activation with

$$h' = h + \alpha \sigma_v v, \quad \alpha \in \{-4, -2, -1, 0, 1, 2, 4\}, \quad (2)$$

where σ_v is the train+validation standard deviation of projections onto v . We sample one boundary from each of 256 test traces, balanced between valid-so-far and invalid-so-far. The outcome is the difference between length-normalized teacher-forced log probabilities of INCORRECT and CORRECT. Twenty matched random unit directions orthogonal to v serve as controls at $\alpha = \pm 4$.

4 Results

4.1 Invalidity is decodable, but first-error timing is harder

Figure 2 shows a rapid rise in decoding performance after the embedding state and stable AUROC above 0.82 through most layers. The validation-selected index obtains test AUROC 0.866 (95% CI [0.849, 0.884]), average precision 0.843 ([0.813, 0.870]), and step F1 0.743 ([0.716, 0.769]). Restricted to erroneous traces, it distinguishes pre-error from error-and-later steps with AUROC 0.873 and average precision 0.925.

The hidden-state score exceeds current-step TF-IDF (AUROC 0.733, Process F1 0.254), position (0.730, 0.054), and the shuffled-label hidden-state control (0.518, 0.179). These controls rule out simple position and current-step lexical explanations, although the present artifacts lack paired confidence intervals for probe-minus-control differences.

Exact timing is substantially weaker. Across erroneous traces, 29.0% are localized exactly, 59.3% within one step, and 72.7% within two. The detector crosses somewhere on 88.7% of erroneous traces but falsely crosses on 38.8% of fully correct traces. Complete-trace exact accuracy is 0.404; Process F1 is 0.393 (95% CI [0.351, 0.436]). Among detected erroneous traces, the mean predicted-minus-human error is +0.59 steps (95% CI [+0.36, +0.80]): the signal is systematically late rather than a sharp internal alarm at the annotated boundary.




Figure 3: Cross-source transfer and causal intervention. Rows are probe-training sources and columns are held-out test sources. Ranking transfers consistently, while threshold-dependent first-error F1 does not. Right: mean paired verdict changes with 95% whole-trace bootstrap intervals; gray points are matched random orthogonal directions at the extreme doses. No learned dose differs reliably from zero.

4.2 Ranking transfers better than localization thresholds

Every off-diagonal transfer AUROC in Figure 3 exceeds 0.739. Mean off-diagonal AUROC is 0.805, compared with 0.845 on the diagonal. The sources therefore share part of the ranking signal. Thresholded localization is less stable: mean off-diagonal Process F1 is 0.261 versus 0.336 on the diagonal, with individual transfers ranging from 0.055 to 0.371. Thus a source-trained direction often orders target-domain boundaries well without transferring calibrated probabilities or the optimal first-crossing threshold.

4.3 The causal assay is inconclusive

Equation 2 produces neither the predicted signed monotonic response nor a nonzero effect at any dose. The mean changes at $\alpha = -4$ and $+4$ are $+0.00040$ ($[-0.00108, 0.00185]$) and -0.00011 ($[-0.00169, 0.00138]$), respectively. Neither extreme beats the distribution of 20 random orthogonal directions.

This null cannot support causal non-use because the unmodified output assay is not informative. All 256 baseline margins favor INCORRECT; recall is 1, specificity is 0, and AUROC is 0.283. Valid boundaries even have a larger mean incorrect-minus-correct margin (5.860) than invalid boundaries (5.625). We can conclude only that this additive intervention, at index 23 and the marked boundary, did not affect this degenerate teacher-forced verdict score.

5 Discussion and limitations

The probe separates prefixes around human-annotated errors and transfers in ranking, but it often fires late and is not shown to drive a correctness judgment. It may encode generator identity, trace-wide difficulty, final-answer outcome, or features richer than TF-IDF. The persistent target may reward accumulated corruption, and the artificial marker may create a summary state. We test one 1.5B model, dataset split, token location, layer, and additive direction. The transfer cells and probe-control gaps also lack paired intervals. A registered follow-up uses within-trace centered probes, a joint surface/metadata control, natural step-end states, a counterbalanced verdict, and a gradient positive control. Decodability establishes accessible information; timing and causal use require separate tests with validated assays.

Responsible use. This study introduces neither human-subject data nor a deployed verifier. Internal error signals could support oversight, but false alarms, delayed detection, distribution shift, and adversarial traces make this probe unsuitable for grading or certifying high-stakes calculations. We report controls and negative results and recommend human or formal verification for consequential uses.

References

- Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219, 2022. doi: 10.1162/coli_a_00422.
- Leonardo Bertolazzi, Philipp Mondorf, Barbara Plank, and Raffaella Bernardi. The validation gap: A mechanistic analysis of how language models compute arithmetic but fail to validate it. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 29387–29424. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.emnlp-main.1495. URL <https://aclanthology.org/2025.emnlp-main.1495/>.
- Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2212.03827>.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021. URL <https://arxiv.org/abs/2110.14168>.
- Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav Goldberg. Amnesic probing: Behavioral explanation with amnesic counterfactuals. *Transactions of the Association for Computational Linguistics*, 9:160–175, 2021. doi: 10.1162/tac1_a_00359. URL <https://aclanthology.org/2021.tac1-1.10/>.
- Bofei Gao, Feifan Song, Zhe Yang, Zefan Cai, Yibo Miao, Qingxiu Dong, Lei Li, Chenghao Ma, Liang Chen, Runxin Xu, Zhengyang Tang, Benyou Wang, Daoguang Zan, Shanghaoran Quan, Ge Zhang, Lei Sha, Yichang Zhang, Xuancheng Ren, Tianyu Liu, and Baobao Chang. Omni-MATH: A universal olympiad-level mathematics benchmark for large language models. *arXiv preprint arXiv:2410.07985*, 2024. URL <https://arxiv.org/abs/2410.07985>.
- Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, Junhao Shen, Jinyi Hu, Xu Han, Yujie Huang, Yuxiang Zhang, Jie Liu, Lei Qi, Zhiyuan Liu, and Maosong Sun. Olympiad-Bench: A challenging benchmark for promoting AGI with olympiad-level bilingual multimodal scientific problems. *arXiv preprint arXiv:2402.14008*, 2024. URL <https://arxiv.org/abs/2402.14008>.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. *arXiv preprint arXiv:2103.03874*, 2021. URL <https://arxiv.org/abs/2103.03874>.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*, pages 2733–2743. Association for Computational Linguistics, 2019. doi: 10.18653/v1/D19-1275. URL <https://aclanthology.org/D19-1275/>.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022. URL <https://arxiv.org/abs/2207.05221>.
- Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems*, volume 36, 2023a. URL <https://arxiv.org/abs/2306.03341>.
- Xiaoyuan Li, Wenjie Wang, Moxin Li, Junrong Guo, Yang Zhang, and Fuli Feng. Evaluating mathematical reasoning of large language models: A focus on error identification and correction. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 11316–11360. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.findings-acl.673. URL <https://aclanthology.org/2024.findings-acl.673/>.

- 222 Yifei Li, Zeqi Lin, Shizhuo Zhang, Qiang Fu, Bei Chen, Jian-Guang Lou, and Weizhu Chen.
223 Making language models better reasoners with step-aware verifier. In *Proceedings of the 61st*
224 *Annual Meeting of the Association for Computational Linguistics*, pages 5315–5333. Association
225 for Computational Linguistics, 2023b. doi: 10.18653/v1/2023.acl-long.291. URL <https://aclanthology.org/2023.acl-long.291/>.
226
- 227 Zhixiang Liang, Beichen Huang, Zheng Wang, and Minjia Zhang. Hidden states as early sig-
228 nals: Step-level trace evaluation and pruning for efficient test-time scaling. In *Findings of*
229 *the Association for Computational Linguistics: ACL 2026*, pages 26800–26813. Association
230 for Computational Linguistics, 2026. doi: 10.18653/v1/2026.findings-acl.1336. URL <https://aclanthology.org/2026.findings-acl.1336/>.
231
- 232 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
233 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. *arXiv preprint*
234 *arXiv:2305.20050*, 2023. URL <https://arxiv.org/abs/2305.20050>.
- 235 Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual
236 associations in GPT. In *Advances in Neural Information Processing Systems*, volume 35, pages
237 17359–17372, 2022. URL <https://arxiv.org/abs/2202.05262>.
- 238 Jingwei Ni, Ekaterina Fadeeva, Tianyi Wu, Mubashara Akhtar, Jiaheng Zhang, Elliott Ash, Markus
239 Leippold, Timothy Baldwin, See-Kiong Ng, Artem Shelmanov, and Mrinmaya Sachan. ReProbe:
240 Efficient test-time scaling of multi-step reasoning by probing internal states of large language
241 models. In *Proceedings of the 64th Annual Meeting of the Association for Computational Lin-*
242 *guistics*, pages 11667–11689. Association for Computational Linguistics, 2026. doi: 10.18653/
243 v1/2026.acl-long.536. URL <https://aclanthology.org/2026.acl-long.536/>.
- 244 Kv Aditya Srivatsa, Kaushal Kumar Maurya, and Ekaterina Kochmar. LLMs cannot spot math
245 errors, even when allowed to peek into the solution. In *Proceedings of the 2025 Confer-*
246 *ence on Empirical Methods in Natural Language Processing*, pages 10914–10928. Association
247 for Computational Linguistics, 2025. doi: 10.18653/v1/2025.emnlp-main.553. URL <https://aclanthology.org/2025.emnlp-main.553/>.
248
- 249 Yucheng Sun, Alessandro Stolfo, and Mrinmaya Sachan. Probing for arithmetic errors in language
250 models. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language*
251 *Processing*, pages 8111–8128. Association for Computational Linguistics, 2025. doi: 10.18653/
252 v1/2025.emnlp-main.411. URL <https://aclanthology.org/2025.emnlp-main.411/>.
- 253 Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udell, Juan J. Vazquez, Ulisse Mini,
254 and Monte MacDiarmid. Steering language models with activation engineering. *arXiv preprint*
255 *arXiv:2308.10248*, 2023. URL <https://arxiv.org/abs/2308.10248>.
- 256 An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu,
257 Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu,
258 Xingzhang Ren, and Zhenru Zhang. Qwen2.5-Math technical report: Toward mathematical
259 expert model via self-improvement. *arXiv preprint arXiv:2409.12122*, 2024. URL <https://arxiv.org/abs/2409.12122>.
260
- 261 Chujie Zheng, Zhenru Zhang, Beichen Zhang, Runji Lin, Keming Lu, Bowen Yu, Dayiheng
262 Liu, Jingren Zhou, and Junyang Lin. ProcessBench: Identifying process errors in mathemat-
263 ical reasoning. In *Proceedings of the 63rd Annual Meeting of the Association for Compu-*
264 *tational Linguistics*, pages 987–1011. Association for Computational Linguistics, 2025. doi:
265 10.18653/v1/2025.acl-long.50. URL <https://aclanthology.org/2025.acl-long.50/>.