
Decodability is not localization: auditing a hidden-state error probe

Anonymous Author(s)

Affiliation

Address

email

Abstract

A linear probe that decodes “this reasoning has gone wrong” from hidden states is often read as a discovery about the model. We check how much of that reading holds up under controls. In Qwen2.5-Math-1.5B-Instruct reading 3,360 ProcessBench traces, a probe on the selected layer predicts the persistent *invalid-so-far* target with held-out AUROC 0.866 [0.849, 0.884], and adding the hidden representation to an equally tuned text-and-metadata model raises AUROC by 0.0587 [0.0414, 0.0745]. Three stronger claims fail. The score ranks boundaries within an erroneous trace almost perfectly (mean within-trace AUROC 0.968) but its threshold names the first wrong step in 28.9% of traces, fires 0.59 steps late on average, and fires on 38.8% of fully correct traces. Ranking transfers across the four ProcessBench sources (off-diagonal AUROC ≥ 0.739) while the thresholded localization metric does not (0.055 to 0.371). The causal test never runs: the model’s own error verdict is below chance on the same boundaries (AUROC 0.342), so the steering results are uninterpretable rather than null. None of the three failures is visible in the pooled AUROC that a probe result usually reports, and we name the check that catches each one.

1 Introduction

Interpretability is increasingly used as a discovery tool. A probe decodes some property from hidden states, and the reported finding is a fact about the model: it represents this, it tracks that. “The model represents whether its own reasoning has become invalid” is a claim of that shape, and it is an appealing one. It promises a cheap step-level verifier and a story about self-monitoring.

A probe licenses less. It shows that a label is linearly decodable from a representation, under a chosen target, split, layer, and threshold. Four things do not follow, and each is a separate empirical question:

1. **Incremental information.** The label may be equally decodable from the visible text and from trace metadata, in which case the hidden state adds nothing.
2. **Localization.** A score that separates pre-error from post-error boundaries need not cross its threshold *at* the error. With a persistent label, every boundary after the error is positive, so accumulated evidence and position are enough to score well.
3. **Portability of the decision rule.** Ranking and thresholded detection can transfer at different rates across data sources.
4. **Causal use.** Decodability does not imply use (Hewitt and Liang, 2019; Belinkov, 2022; Elazar et al., 2021). Testing use requires a behavioral readout that works before it is perturbed.

We ran all four checks on one system: Qwen2.5-Math-1.5B-Instruct (Yang et al., 2024) reading 3,360 human-annotated mathematical reasoning traces from ProcessBench (Zheng et al., 2025). Check 1 passes. Checks 2 and 3 fail in specific ways. Check 4 cannot be run at all, because the

model’s own verdict about its reasoning is worse than chance on the boundaries we would have intervened on, and we report that as a failed precondition rather than as a null effect. The model here is a test case. What matters is that all four failures are compatible with the headline number. Pooled AUROC of 0.866 is reported the same way whether or not a text baseline reaches 0.810, whether the threshold fires early or late, and whether the model’s behavior responds to the direction. An interpretability result read as a discovery needs the checks that separate those cases. Here every check runs on cached activations and stored predictions, on CPU, and together they change what the study is allowed to say.

2 Related work

Process supervision scores intermediate reasoning steps rather than final answers (Cobbe et al., 2021; Lightman et al., 2023; Li et al., 2023). ProcessBench annotates the earliest erroneous step in traces drawn from GSM8K, MATH, OlympiadBench, and Omni-MATH (Zheng et al., 2025; Hendrycks et al., 2021; He et al., 2024; Gao et al., 2024).

Several recent studies score reasoning from internal states. Sun et al. (2025) probe arithmetic errors and re-prompt selectively on the result. ReProbe uses hidden-state step scoring for test-time scaling and STEP prunes traces with a hidden-state scorer (Ni et al., 2026; Liang et al., 2026). Alvarez and Baheri (2026) study first-error detection with hidden-state transport geometry across models and datasets. Yuan et al. (2026) separate diagnostic error awareness from causal use across several models. Our scope is narrower than any of these and our target is different: we hold the model and benchmark fixed and vary the evidential standard, so we measure the gap between a decoding result and a localization or causal result on one set of predictions instead of arguing for it.

3 Setup

Data and split. Each ProcessBench trace has a problem, a generator identity, a source, written steps, final-answer correctness, and a human first-error index e_i , with $e_i = -1$ for a fully correct trace. Traces sharing a normalized problem statement go to the same partition; seed 42 gives 60/20/20 train, validation, and test partitions. Of 3,400 attempted traces we retain 3,360 and 24,909 step boundaries; the 40 exclusions exceed the 2,048-token context limit. The test partition holds 669 traces (432 erroneous, 237 correct) and 4,985 boundaries. The model reads each fixed trace in one causal forward pass and does not generate it. Appendix A gives the full accounting.

Target and probe. At boundary k of trace i the target is $y_{ik} = \mathbf{1}[e_i \geq 0 \wedge k \geq e_i]$: the prefix is invalid so far. It marks the first invalid step and every later boundary, and it does not say whether the current step is locally wrong. Extraction records the 1,536-dimensional residual-stream state at each step boundary at 29 hidden-state indices. At each index we fit a standardized L2 logistic probe with class balancing over $C \in \{0.01, 0.1, 1, 10\}$. Validation AUROC selects index 23 and $C = 0.01$; validation Process F1 selects threshold 0.645 on a grid from 0.05 to 0.95. The probe is refit on train plus validation and evaluated once on test. The predicted first error is the first threshold crossing, or -1 if the score never crosses. Intervals are whole-trace bootstraps (1,000 draws for the probe and the nested comparison, 2,000 for the control audit).

Controls. Every nuisance model gets the probe’s selection budget: the same C grid, a validation-selected threshold, a refit on train plus validation, and inverse boundary-count weights so each trace contributes equal training loss. Prefix TF-IDF sees the problem and reasoning text through the current boundary. Structural metadata sees source, generator, absolute and fractional position, trace length, and token count. A post-hoc semantic control encodes the same visible prefix with mean-pooled all-MiniLM-L6-v2 embeddings under the same fitting procedure; its saved record does not pin an encoder revision or keep paired bootstrap draws, so we report point estimates. Final-answer correctness is a benchmark annotation that no online detector would have, so we label models that use it diagnostic and never treat them as baselines.

Post-hoc analyses. The temporal, boundary-location, transition, and semantic-baseline analyses were designed after the primary test result was known. We mark them post-hoc throughout. Their intervals describe resampling uncertainty inside the completed pipeline, not sensitivity to a new split or selection procedure.




Figure 1: Left: held-out boundary AUROC and thresholded first-error F1 at each hidden-state index. The gap is present at every index, including the validation-selected one. Center: mean probe score aligned to the annotated first error, with a 95% whole-trace bootstrap band. At the onset the mean score is 0.559, below the operating threshold of 0.645; it crosses one step later. Right: the 16 source-pair evaluations. AUROC stays above 0.739 off the diagonal while first-error F1 falls as low as 0.055.

Table 1: Held-out comparison on 669 test traces. Every control receives the probe’s selection budget and trace-equal training weights. Rows marked $\dagger$ use final-answer correctness, a benchmark annotation unavailable to an online detector; they are diagnostic, not baselines.

Predictor	AUROC	Proc. F1	Error exact	Corr. rej.
Hidden state (selected layer)	.866	.393	.289	.612
Prefix TF-IDF	.751	.274	.192	.477
Prefix MiniLM embeddings (post-hoc)	.753	.268	—	—
Structural metadata	.776	.168	.148	.194
Text + metadata	.810	.210	.169	.278
Metadata + final answer $\dagger$	.854	.262	.164	.646
Text + metadata + final answer $\dagger$	.874	.294	.176	.895

In the separately fit nested comparison (Appendix C), adding the selected-layer hidden representation raises AUROC by .0587 [.0414, .0745] over text plus metadata and by .0348 [.0242, .0451] over text, metadata, and final answer.

89 4 Results

90 4.1 Check 1: hidden states add information the visible text does not

91 The selected probe reaches test AUROC 0.866 [0.849, 0.884], average precision 0.843
 92 [0.813, 0.870], and step F1 0.743 [0.716, 0.769]. The equally tuned controls are not far behind:
 93 0.751 for prefix TF-IDF, 0.753 for the MiniLM prefix encoder, 0.776 for structural metadata, and
 94 0.810 for text and metadata jointly (Table 1). A model given metadata plus final-answer correctness
 95 reaches 0.854, and one given text, metadata, and final-answer correctness reaches 0.874, above the
 96 probe. That field is an annotation, not a signal, but it shows how much of this target a benchmark’s
 97 own bookkeeping predicts.

98 The nested comparison is the direct test. Let N be prefix text plus metadata and $N + H$ the
 99 same features plus the full selected-layer hidden representation. Adding H improves AUROC by
 100 0.0587 [0.0414, 0.0745], log loss by -0.0798 [-0.1038 , -0.0548], and Process F1 by $+0.1866$
 101 [0.1397, 0.2340]. Adding H on top of the diagnostic feature set that already contains final-answer
 102 correctness still improves AUROC by 0.0348 [0.0242, 0.0451]. Both intervals exclude zero on this
 103 split. Check 1 passes, though the gain is modest: a tuned bag of words and six pieces of trace
 104 metadata recover most of what the probe recovers.

105 4.2 Check 2: the same score does not localize the first error

106 Within an erroneous trace, the probe orders boundaries almost perfectly: mean within-trace AUROC
 107 is 0.968 [0.958, 0.976] over the 380 test traces whose first error is not at step 0. Any score that rises
 108 with step index would rank well on that measure, so it is not evidence about the representation. It
 109 is the ceiling that the decision rule built from the same scores fails to reach. That rule names the
 110 exact first-error step in 28.9% of erroneous traces, lands within one step in 59.3%, and within two

in 72.7%. It crosses at all on 88.7% of erroneous traces and also crosses on 38.8% of fully correct ones. Among detected traces the mean signed error is +0.59 steps [+0.36, +0.80]: the crossing comes after the annotation, not before it.

Figure 1 (center) shows why. The mean score rises from 0.341 one step before the annotated error to 0.559 at the onset, 0.680 one step later, and 0.731 two steps later. At the onset the average trace is still below the operating threshold. The signal is a ramp, not an event, and a fixed threshold on a ramp fires late.

The alignment to the annotation is nonetheless real. Circularly shifting each frozen score trajectory 5,000 times, which preserves its values and destroys its alignment, gives a null exact rate of 0.164 [0.137, 0.192] against the observed 0.289 ($p = 0.0002$), and a null Process F1 of 0.259 against the observed 0.393 ($p = 0.0002$). Comparing the score jump at a noninitial first error with the jump at a transition in a correct trace matched on source, generator, relative position, trace length, and token count gives 0.257 against 0.113, a paired difference of 0.144 [0.096, 0.193]. Correct traces drift upward too, which is the same ramp seen without an error to explain it.

Two robustness checks are flat. Choosing thresholds per length bucket moves Process F1 from 0.393 to 0.377 and correct-trace rejection from 0.612 to 0.540, leaving exact localization unchanged. Reading states at the last natural token of each step rather than at the artificial step marker changes AUROC by +0.002 [−0.007, +0.011] and Process F1 by +0.026 [−0.020, +0.075], so the result is not an artifact of the boundary token.

4.3 Check 3: ranking transfers, the operating point does not

Training a probe on one ProcessBench source and testing on another gives mean off-diagonal AUROC 0.805 against mean diagonal 0.845, with every off-diagonal value above 0.739. The thresholded metric moves much more: mean off-diagonal Process F1 is 0.261 against 0.336 on the diagonal, with off-diagonal values from 0.055 to 0.371 (Figure 1, right). The worst cell, GSM8K to OlympiadBench, keeps AUROC 0.753 while Process F1 falls to 0.055: the threshold fitted on short grade-school traces fires on 97% of correct OlympiadBench traces. These are four sources inside one benchmark construction, not four datasets, and they say nothing about another model family or size.

4.4 Check 4: the causal test does not run

To ask whether the model uses the decoded direction we need a behavioral readout that distinguishes valid from invalid prefixes before any intervention. We asked the model whether its reasoning contains an error and scored the difference between the conditional probabilities of “Yes” and “No”. On 256 balanced boundaries this readout has AUROC 0.342 and specificity 0.203: worse than chance at the task the intervention would have measured. An earlier single-token CORRECT versus INCORRECT readout was worse, at AUROC 0.283 and specificity zero.

The pipeline gates the causal analysis on that readout clearing AUROC 0.5 with nonzero specificity, so we report the intervention as an assay diagnostic. We moved the boundary state along the probe direction at $\alpha \in \{-4, -2, -1, 0, 1, 2, 4\}$ projection standard deviations and also along 20 random orthogonal directions at the extreme doses; every learned-direction interval includes zero, with a smallest detectable effect of 0.0029 at 80% power. None of that is evidence about the model. A flat dose-response curve on a broken instrument is uninformative, and reporting it as a null result would be the mistake the gate exists to prevent.

An exploratory positive. A probe on the onset difference $h_{i,e_i} - h_{i,e_i-1}$ separates error onsets from matched correct transitions with AUROC 0.769 [0.713, 0.820] on 380 test pairs, above position (0.631), destination-step TF-IDF (0.671), and shuffled labels (0.585). The matching is imperfect: one placebo trace supports up to 46 pairs, and standardized differences are 0.146 for trace length and 0.285 for token count. The test split had been inspected before this analysis was conceived. We report it as exploratory and draw no mechanism claim from it.

5 What this changes about reading probe results

Four things generalize past this model, and each is cheap to check.

Report the decision rule, not the ranking. Within-trace AUROC of 0.968 and exact localization of 28.9% come from the same predictions. If the discovery claim is that a model knows *when* its reasoning went wrong, AUROC is the wrong summary, because it never touches the threshold that would answer the question.

165 **Give the nuisance model the probe’s tuning budget.** Text plus metadata reaches 0.810 against
166 the probe’s 0.866. Measured against prefix TF-IDF alone the gap is 0.115, about twice the 0.0587
167 that survives the joint nuisance model, and measured against no control at all the reader sees 0.866
168 against chance.

169 **Watch for benchmark bookkeeping in the feature set.** Adding final-answer correctness, a field
170 that sits in the ProcessBench metadata, lifts the nuisance model to AUROC 0.874, past the probe,
171 while being unavailable to any detector running before the answer exists. Fields like this are easy to
172 include by accident and hard to notice afterward.

173 **Gate causal claims on assay validity.** We could have run the steering experiment, found flat curves,
174 and written that the direction is decoded but unused. The readout was below chance, so that sentence
175 would have had no support. Checking the readout first left us with a gap in the evidence instead of
176 a false null.

177 6 Limitations

178 One model, one benchmark construction, one split seed. The nested gain in Table 1 is conditional on
179 that partition; we have not repeated the grouped split, layer and C selection, and threshold selection
180 across seeds, which is the first thing we would run next and needs no new extraction. The persistent
181 target labels every boundary after the annotated error, so pooled metrics reward accumulated
182 evidence. The MiniLM control is one unpaired post-hoc run with an unpinned encoder revision and
183 does not exhaust visible-text representations. Calibration is not evaluated as a separate question;
184 we report log loss and Brier score but no reliability analysis. The transition result is exploratory
185 for the reasons given above, and the planned inverse-reuse and one-to-one matching refits have no
186 completed artifact. Counterfactual activation patching produced no result: its corrections were never
187 independently verified and it inherits the same failed behavioral gate.

188 Responsible use statement

189 The probe studied here is not a grader. At its selected threshold it flags 38.8% of fully correct so-
190 lutions as containing an error and names the wrong step in most of the traces where an error exists,
191 so using it for scoring, certification, or unattended feedback on student or user work would produce
192 frequent and unevenly distributed false accusations; held-out correct-trace rejection already varies
193 from 0.463 on OlympiadBench to 0.786 on GSM8K. The supported use is as a research diagnostic
194 alongside human or formal verification. The same point applies past this probe: a score presented as
195 evidence that a model “knows” it is wrong can be used to justify automation that the underlying evi-
196 dence does not support, and the controls in Table 1 and the checks in Section 5 exist to make that gap
197 visible. The data are model-generated mathematical solutions from a public benchmark and contain
198 no personal information. The analysis package, configuration, frozen prediction artifacts, and figure
199 code that reproduce every number here are included in the anonymized supplement; extraction needs
200 one A100-class GPU and the remaining analyses run on CPU.

201 References

- 202 Tyler Alvarez and Ali Baheri. Where does reasoning break? step-level hallucination detection via
203 hidden-state transport geometry. *arXiv preprint arXiv:2605.13772*, 2026. URL <https://arxiv.org/abs/2605.13772>.
- 204
- 205 Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational*
206 *Linguistics*, 48(1):207–219, 2022. doi: 10.1162/coli_a_00422.
- 207 Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser,
208 Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John
209 Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*,
210 2021. URL <https://arxiv.org/abs/2110.14168>.
- 211 Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav Goldberg. Amnesic probing: Behavioral
212 explanation with amnesic counterfactuals. *Transactions of the Association for Computational*
213 *Linguistics*, 9:160–175, 2021. doi: 10.1162/tac1_a_00359. URL [https://aclanthology.org/](https://aclanthology.org/2021.tac1-1.10/)
214 [2021.tac1-1.10/](https://aclanthology.org/2021.tac1-1.10/).
- 215 Bofei Gao, Feifan Song, Zhe Yang, Zefan Cai, Yibo Miao, Qingxiu Dong, Lei Li, Chenghao Ma,
216 Liang Chen, Runxin Xu, Zhengyang Tang, Benyou Wang, Daoguang Zan, Shanghaoran Quan,
217 Ge Zhang, Lei Sha, Yichang Zhang, Xuancheng Ren, Tianyu Liu, and Baobao Chang. Omni-

- MATH: A universal olympiad-level mathematics benchmark for large language models. *arXiv preprint arXiv:2410.07985*, 2024. URL <https://arxiv.org/abs/2410.07985>.
- Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, Junhao Shen, Jinyi Hu, Xu Han, Yujie Huang, Yuxiang Zhang, Jie Liu, Lei Qi, Zhiyuan Liu, and Maosong Sun. Olympiad-Bench: A challenging benchmark for promoting AGI with olympiad-level bilingual multimodal scientific problems. *arXiv preprint arXiv:2402.14008*, 2024. URL <https://arxiv.org/abs/2402.14008>.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. *arXiv preprint arXiv:2103.03874*, 2021. URL <https://arxiv.org/abs/2103.03874>.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*, pages 2733–2743. Association for Computational Linguistics, 2019. doi: 10.18653/v1/D19-1275. URL <https://aclanthology.org/D19-1275/>.
- Yifei Li, Zeqi Lin, Shizhuo Zhang, Qiang Fu, Bei Chen, Jian-Guang Lou, and Weizhu Chen. Making language models better reasoners with step-aware verifier. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, pages 5315–5333. Association for Computational Linguistics, 2023. doi: 10.18653/v1/2023.acl-long.291. URL <https://aclanthology.org/2023.acl-long.291/>.
- Zhixiang Liang, Beichen Huang, Zheng Wang, and Minjia Zhang. Hidden states as early signals: Step-level trace evaluation and pruning for efficient test-time scaling. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 26800–26813. Association for Computational Linguistics, 2026. doi: 10.18653/v1/2026.findings-acl.1336. URL <https://aclanthology.org/2026.findings-acl.1336/>.
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. *arXiv preprint arXiv:2305.20050*, 2023. URL <https://arxiv.org/abs/2305.20050>.
- Jingwei Ni, Ekaterina Fadeeva, Tianyi Wu, Mubashara Akhtar, Jiaheng Zhang, Elliott Ash, Markus Leippold, Timothy Baldwin, See-Kiong Ng, Artem Shelmanov, and Mrinmaya Sachan. ReProbe: Efficient test-time scaling of multi-step reasoning by probing internal states of large language models. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics*, pages 11667–11689. Association for Computational Linguistics, 2026. doi: 10.18653/v1/2026.acl-long.536. URL <https://aclanthology.org/2026.acl-long.536/>.
- Yucheng Sun, Alessandro Stolfo, and Mrinmaya Sachan. Probing for arithmetic errors in language models. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 8111–8128. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.emnlp-main.411. URL <https://aclanthology.org/2025.emnlp-main.411/>.
- An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu, Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu, Xingzhang Ren, and Zhenru Zhang. Qwen2.5-Math technical report: Toward mathematical expert model via self-improvement. *arXiv preprint arXiv:2409.12122*, 2024. URL <https://arxiv.org/abs/2409.12122>.
- Aojie Yuan, Zhiyuan Julian Su, Haiyue Zhang, Yi Nian, and Yue Zhao. Hidden error awareness in chain-of-thought reasoning: The signal is diagnostic, not causal. *arXiv preprint arXiv:2605.09502*, 2026. URL <https://arxiv.org/abs/2605.09502>.
- Chujie Zheng, Zhenru Zhang, Beichen Zhang, Runji Lin, Keming Lu, Bowen Yu, Dayiheng Liu, Jingren Zhou, and Junyang Lin. ProcessBench: Identifying process errors in mathematical reasoning. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics*, pages 987–1011. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.acl-long.50. URL <https://aclanthology.org/2025.acl-long.50/>.

269 A Data accounting

270 Retained boundary counts by source, in the order shown above, are 1,229 / 3,718 / 5,129 / 4,868
 271 for train, 414 / 1,303 / 1,708 / 1,555 for validation, and 439 / 1,387 / 1,661 / 1,498 for test. The
 272 attempted total was 25,697 boundaries and the retained total is 24,909.

Table 2: Attempted and retained traces and boundaries by frozen partition.

Partition	Attempted	Retained	Excluded	Boundaries	Error traces	Correct traces
Train	2,041	2,014	27	14,944	1,311	703
Validation	681	677	4	4,980	444	233
Test	678	669	9	4,985	432	237
All	3,400	3,360	40	24,909	2,187	1,173

Table 3: Attempted/retained/excluded traces by source and frozen partition.

Partition	GSM8K	MATH	OlympiadBench	Omni-MATH
Train	236/236/0	597/592/5	602/594/8	606/592/14
Validation	80/80/0	198/198/0	204/202/2	199/197/2
Test	84/84/0	205/204/1	194/188/6	195/193/2

B Full metric table and subgroup audit

The selected probe’s held-out metrics are AUROC 0.866 [0.849, 0.884], average precision 0.843 [0.813, 0.870], step F1 0.743 [0.716, 0.769], error exactness 0.289, correct-trace rejection 0.612, Process F1 0.393, and complete-trace accuracy 0.404. Weighting traces equally at evaluation rather than by boundary count changes AUROC by +0.0013 [−0.0069, 0.0103], average precision by −0.0052, and step F1 by −0.0087, all with intervals covering zero.

Table 4: Trace-level outcomes by ProcessBench source, with 95% whole-trace bootstrap intervals. These are descriptive subgroup results, not independent hypothesis tests.

Source	Complete accuracy	Exact error	Correct rejection
GSM8K	.512 [.405, .619]	.238 [.116, .375]	.786 [.658, .902]
MATH	.500 [.426, .569]	.347 [.264, .435]	.723 [.620, .818]
OlympiadBench	.298 [.234, .362]	.207 [.138, .282]	.463 [.340, .585]
Omni-MATH	.358 [.290, .425]	.324 [.253, .401]	.467 [.319, .614]

Complete-trace accuracy is 0.292 [0.245, 0.342] among traces with an incorrect final answer and 0.506 [0.451, 0.554] among traces with a correct one. Across generators contributing at least 20 held-out traces it ranges from 0.316 to 0.621.

C Nested comparison and intervention records

D Reproduction record

The package exposes one configuration file and one command group:

```

uv sync --extra dev
uv run math-error --config configs/project.yaml validate-config
uv run math-error --config configs/project.yaml run-all
uv run math-error --config configs/project.yaml analyze
uv run math-error --config configs/project.yaml fit-conditional
uv run math-error --config configs/project.yaml fit-contextual-baseline
uv run math-error --config configs/project.yaml fit-transition
uv run math-error --config configs/project.yaml transition-diagnostics
uv run math-error --config configs/project.yaml analyze-boundaries
uv run pytest && uv run ruff check src tests

```

The configuration records seed 42, model and dataset identities, the 2,048-token limit, selection grids, bootstrap counts, and artifact paths. Extraction ran on one NVIDIA A100 in bfloat16 at batch size 16 over single causal passes; the extraction record fixes the model identity, dtype, token limit, and dataset SHA-256. Probe fitting and every audit in this paper run on CPU, and the audit record stores Python, NumPy, pandas, SciPy, scikit-learn, matplotlib, and platform versions. Dataset and model use follows the licenses of the published sources.

Table 5: Nested feature-set comparison. N is prefix TF-IDF plus structural metadata; O adds final-answer correctness; H is the full selected-layer hidden representation. All conditions share the partitions, the C grid, trace-equal training weights, and the refit protocol.

Condition	AUROC	Log loss	Proc. F1	Δ AUROC vs. nuisance only
N	.811	.533	.211	
$N + H$	.869	.454	.397	+.0587 [.0414, .0745]
O	.874	.440	.284	
$O + H$	.909	.385	.444	+.0348 [.0242, .0451]
H alone	.868	.457	.397	




Figure 2: Paired change in the model’s verdict score under the learned-direction intervention at each dose, with 95% intervals. Every interval covers zero. Because the unperturbed readout has AUROC 0.342, this figure records an assay diagnostic and carries no evidence about causal use.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and Introduction name the single model and benchmark, distinguish the persistent prefix target from onset detection, and state the negative and post-hoc results.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Section 6 discusses one-model and one-split scope, shortcut sensitivity, the post-hoc status of several analyses, the untested calibration question, and the invalid causal assay.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper contains no theoretical results or proof claims.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 3 and Appendix D specify the split, target, selection protocol, seeds, locked package versions, dataset identity, and exact reproduction commands. The joint shortcut controls share the hidden probe’s validation budget, and the transition-matching protocol is stated with its reuse diagnostics.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: An anonymized code supplement includes the pipelines, tests, configurations, and instructions; the upstream dataset and model are publicly accessible.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.

- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 3 reports the data partitions, exclusions, model and activation location, probe family, hyperparameter grid, threshold selection, controls, and estimands.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The main results use 95% whole-trace bootstrap intervals, paired control differences, and a two-way clustered bootstrap for reused transition controls.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix D records the GPU model, precision, and batch size for extraction, and states that probe fitting and every audit reported here run on CPU. Exact wall-clock logs for the original extraction were not preserved, so the paper gives the hardware and settings rather than inventing a runtime figure.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work analyzes publicly released model and benchmark artifacts, makes bounded claims, preserves submission anonymity, and reports failures and limitations.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The responsible-use discussion identifies research-diagnostic value and the concrete risks of false alarms, late detections, and distribution shift in grading uses.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases no new pretrained model, scraped dataset, or other asset with a high misuse risk; it releases analysis code for public research artifacts.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The paper cites the ProcessBench benchmark and the Qwen model, and Appendix D states that dataset and model use follows the licenses published by those sources, whose constituent benchmarks retain their original terms.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The anonymous supplement documents the analysis code, resolved configurations, artifact schemas, commands, tests, and known limitations.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.

- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: This study performs no crowdsourcing and recruits no human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: This study performs no human-subject research.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: The analyzed language model is the experimental subject, not a non-standard agent used to generate labels, hypotheses, or core methodological decisions. Agent assistance contributed to code, analysis, and manuscript editing under author direction; the authors verified all reported numbers against repository artifacts and remain responsible for correctness and originality.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.

641
642

- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.